

// CHEATSHEET

THE AI PROMPT ENGINEERING CHEATSHEET

10 Techniques That Actually Work

A field guide to writing AI prompts that get sharper, more accurate, more useful results from the first try. Save it. Use it. Reach for it next time you open ChatGPT, Claude, or any model.

BY RYU VY

AIToolsDaily24 | By a human, for humans

NEW TECHNIQUES WEEKLY

longvy24.github.io/AI_Blog/index.html

// HOW TO USE THIS

A Field Guide, Not A Textbook.

These ten techniques are the ones I actually use every day. Each one is a small change to how you write your prompts, and each one materially improves what AI gives back.

The format is the same for all ten: what the technique does, the exact prompt template, an example, when to reach for it, and one pro tip from real use. Pages are built to be screenshot-able. Save the ones you use most.

These work on ChatGPT, Claude, Gemini, and any modern AI model. No plugins. No setup. Just the prompt itself.

// THE TEN TECHNIQUES

01	Role-Based Prompting	Give AI a job title.
02	Few-Shot Prompting	Show examples. AI pattern-matches.
03	Chain-of-Thought	Make AI think step by step.
04	Iterative Refinement	Edit your prompt, not the output.
05	Delimiters	Wrap pasted content in tags.
06	Instructional Prompting	Start every prompt with a verb.
07	Negative Constraints	Tell AI what NOT to do.
08	Prompt Chaining	Break tasks into sequential steps.
09	Retrieval Augmentation	Paste your source. Use only that.
10	Emotion Prompting	Add stakes. AI responds to context.

01

ROLE-BASED PROMPTING

Tell the model who to be before you tell it what to do.

// WHAT IT DOES

Assigning the model a specific role gives it context it would otherwise have to guess at. Tone, vocabulary, expertise level, and perspective all shift instantly. The single highest-leverage one-line change in prompting.

// THE PROMPT

```
Act as a [specific role] with [X years] of experience in [field].
```

```
Then answer:
```

```
[your actual question or task]
```

// EXAMPLE

INPUT

Act as a senior B2B SaaS marketing manager with 10 years of experience. Then answer: How do I write a cold outreach email that does not sound spammy?

OUTPUT

Output reads like advice from someone who has actually written B2B emails: specific subject-line patterns, opener formulas, a structure that mirrors how SaaS sales cycles actually work.

// WHEN TO USE

- > Any domain where expertise matters: finance, legal, design, technical writing, medical.
- > When default output feels generic, surface-level, or wrong-toned.

// PRO TIP

Specificity wins. "Marketing expert" is fine. "Senior B2B SaaS marketing manager with 10 years in enterprise sales" is better. Narrow the role and the output sharpens with it.

02

FEW-SHOT PROMPTING

Show the model the pattern. It will copy it.

// WHAT IT DOES

Models pattern-match. Give them two or three input-output examples before your actual task and the response will mirror the format, tone, and structure of your examples almost exactly. Faster than describing what you want in words.

// THE PROMPT

```
Here are examples.
```

```
Input: [example 1] -> Output: [desired output 1]
```

```
Input: [example 2] -> Output: [desired output 2]
```

```
Input: [example 3] -> Output: [desired output 3]
```

```
Now do: [your actual task]
```

// EXAMPLE

INPUT

Show 3 example email subject lines in a punchy direct style, then ask the model to write a fourth in the same style for a new product launch.

OUTPUT

The fourth subject line lands in the same voice as the three examples: same word count band, same energy, same structure. No vibe drift.

// WHEN TO USE

- > Writing-style matching, custom email templates, structured data extraction, format-sensitive tasks.
- > When describing what you want is harder than showing it.

// PRO TIP

Three examples is the sweet spot. One is too few. Five-plus is diminishing returns. If the output still drifts, your examples are not consistent enough with each other.

03

CHAIN-OF-THOUGHT (COT)

Three words that make the model reason instead of guess.

// WHAT IT DOES

Adding "think step by step" forces the model to work through the problem out loud before answering. Accuracy on math, logic, and multi-step reasoning improves dramatically. The model shows its work, which also makes mistakes easier to catch.

// THE PROMPT

```
[your question]
```

```
Think step by step.
```

// EXAMPLE

INPUT

A meeting room holds 12. Five teams want to meet for an hour, two need two hours back-to-back, one needs 30 minutes. The room is free 9-5. Can it fit? Think step by step.

OUTPUT

Model walks through each booking, adds the durations, checks for conflicts, and answers with the reasoning visible. Mistakes become obvious before they hit you.

// WHEN TO USE

- > Math problems, logic puzzles, scheduling, business decisions.
- > Anything where accuracy matters more than response speed.

// PRO TIP

Combine with "Show your work" or "Explain your reasoning" for even more thorough chains. For Claude specifically, ask it to use a thinking block.

04

ITERATIVE REFINEMENT

Stop editing AI output. Edit your prompt instead.

// WHAT IT DOES

Real AI users do not write a single perfect prompt. They write a rough one, see the output, then add constraints in rounds: tone, then format, then length, then negative constraints. Four iterations beats one shot every time.

// THE PROMPT

```
Round 1: [basic prompt]
Round 2: [+ tone, e.g. "professional, no fluff"]
Round 3: [+ format, e.g. "3 bullets, then a one-line summary"]
Round 4: [+ length, e.g. "under 80 words total"]
Round 5: [+ constraints, e.g. "no preamble, no jargon"]
```

// EXAMPLE

INPUT

Start: "Write me an email." After 4 rounds: "Write a 3-sentence follow-up email to a warm B2B lead. Conversational tone. End with one clear CTA. No preamble."

OUTPUT

Iteration 1 gives a generic email. Iteration 4 gives publish-ready copy. Same model, same conversation, dramatically better output.

// WHEN TO USE

- > Long-form writing, anything that needs to be polished or publish-ready.
- > When the first output is "close but not quite."

// PRO TIP

Save your refined final prompt as a template. You will reuse it. Most of my workhorse prompts started life as iteration 4 of something messy.

05

DELIMITERS

A tiny syntax trick that stops AI from misreading you.

// WHAT IT DOES

When you paste a long article, document, or block of code into a prompt, the model can confuse your instructions with your data. Wrap your content in clear delimiters and the confusion goes away. The output gets noticeably cleaner.

// THE PROMPT

```
Summarize the article below in 3 bullets.
```

```
"""
```

```
[your pasted article goes here]
```

```
"""
```

// EXAMPLE

INPUT

Pasting a 1,500-word product memo and asking for a 3-bullet summary. Without delimiters, the model sometimes summarizes YOUR instruction instead of the content.

OUTPUT

With triple-quote delimiters, you get a clean 3-bullet summary of the actual memo. Every time. The model knows where your instruction ends and the data begins.

// WHEN TO USE

- > Any prompt that pastes in source content: articles, code, emails, transcripts, support tickets.
- > When the model keeps "summarizing the wrong thing."

// PRO TIP

Triple quotes work in most models. Use XML-style tags like `<source> ... </source>` for Claude. Use `###` for code-heavy prompts where quotes might collide with the content.

06

INSTRUCTIONAL PROMPTING

Start every prompt with a verb. Every one.

// WHAT IT DOES

Models are trained to follow commands, not interpret hints. "Help me with my email" is a hint. "Write a 3-sentence email" is a command. The verb-first version gets a dramatically more useful response. This is foundational.

// THE PROMPT

[Verb] [specific output] [for/about/in the style of...]

Examples of verbs to start with:

Summarize. Analyze. Write. Compare. List. Draft. Rewrite. Outline. Critique. Translate.

// EXAMPLE**INPUT**

Bad: "Help me think about my pricing." Good: "Compare three pricing models for a \$20/mo B2B SaaS and list the pros and cons of each in a table."

OUTPUT

Hint version returns a vague pep talk. Verb-first version returns the table you actually wanted. Same question, opposite results.

// WHEN TO USE

- > Every single prompt. This is not a "sometimes" technique.
- > Especially when output keeps coming back vague or wandering.

// PRO TIP

My go-to verbs in priority order: Summarize, Analyze, Write, Compare, List, Rewrite, Critique, Draft, Outline. Pick the sharpest one for the job.

07

NEGATIVE CONSTRAINTS

Tell the model what NOT to do. Everything changes.

// WHAT IT DOES

Telling AI what to avoid is often more powerful than telling it what to do. It removes the most common annoyances (preambles, jargon, recap) without you having to micromanage everything else. The output gets shorter, sharper, and easier to use.

// THE PROMPT

```
[your request]
```

```
Do not [thing 1].
```

```
Do not [thing 2].
```

```
Do not [thing 3].
```

// EXAMPLE

INPUT

Summarize the meeting notes below. Do not use jargon. Do not exceed 5 bullet points. Do not include a preamble or a recap at the end.

OUTPUT

You get five bullets. No "Here is a summary of..." preamble. No "In conclusion..." recap. Just the content. Saves about a minute of cleanup per output.

// WHEN TO USE

- > When the output is technically correct but stylistically wrong.
- > When you find yourself deleting the same kind of fluff every time.

// PRO TIP

My standard "do not"s added to every prompt by default: no preamble, no jargon, no recap at the end. Those three alone fix 80 percent of output annoyances.

08

PROMPT CHAINING

One big prompt is one big mess. Break it into steps.

// WHAT IT DOES

Complex tasks (blog posts, scripts, reports) come out better when broken into a sequence of simple prompts. Each step has one job. Each output becomes the next input. Quality compounds instead of degrading.

// THE PROMPT

```
Step 1: Generate an outline of [topic].
Step 2: Expand each outline section into 2-3 paragraphs.
Step 3: Write the intro using the outline as context.
Step 4: Rewrite the whole piece for [target tone].
```

// EXAMPLE**INPUT**

Writing a 1,200-word blog post. Step 1 asks for outline. Step 2 expands each section separately. Step 3 writes intro. Step 4 polishes tone. Each step is its own prompt.

OUTPUT

Final output is dramatically more coherent and on-topic than asking for the whole post in one shot. Each step gives you a chance to course-correct before the next.

// WHEN TO USE

- > Anything longer than a paragraph: blog posts, scripts, reports, briefs, technical analyses.
- > When one-shot prompts produce something close but uneven.

// PRO TIP

Save each step output before moving to the next. You will often want to revise an upstream step without re-running every step that follows.

09

RETRIEVAL AUGMENTED GENERATION

Paste your source. Tell it to use only that.

// WHAT IT DOES

AI hallucinations happen because the model fills in gaps when it does not actually know. The fix: give it your own source material and tell it to use only that. The model becomes a careful reader instead of a guess machine.

// THE PROMPT

```
Answer only from the text below.  
If the answer is not in the text, say "I do not know."  
  
""  
[your source text]  
""  
  
Question: [your question]
```

// EXAMPLE

INPUT

Pasting a customer support FAQ and asking specific product questions. The model can only answer from the FAQ. No hallucinated features, no made-up policies.

OUTPUT

If the answer is in your source, you get it cleanly. If it is not, the model says so instead of inventing something. This is the difference between AI as a research tool and AI as a liability.

// WHEN TO USE

- > Research, FAQs, legal documents, product specs, internal knowledge bases.
- > Anywhere accuracy is non-negotiable and the source material is available.

// PRO TIP

This is the prompt that turns AI from "confident bullshitter" into "grounded reader." Use it any time you would otherwise have to fact-check the output.

10

EMOTION PROMPTING

Add stakes to the prompt. The model responds.

// WHAT IT DOES

Telling the model "this is important to me" actually changes how it responds. The output gets longer, more thorough, more considered. There is real research behind this; it is not a placebo. Models read context, and effort signals are part of that context.

// THE PROMPT

```
[your request]
```

```
This is very important for my [career/team/project].
```

// EXAMPLE

INPUT

Review this draft cover letter and tell me what to improve. This is very important for my career.

OUTPUT

Same model, same question, with the stakes line: longer response, more specific suggestions, more careful tone. Visibly more effort.

// WHEN TO USE

- > When you want the model to take more time and be more thorough.
- > High-stakes deliverables: cover letters, presentations, important emails.

// PRO TIP

Do not overdo it. "Important for my career" works. "My life depends on this" reads as manipulation and the model often pushes back or refuses. Real stakes only.

// PRINT THIS PAGE

All Ten, One Page.

01 Role-Based

Act as a [role] with [X] years experience.

02 Few-Shot

Input: [X] -> Output: [Y]. Now do: [task].

03 Chain-of-Thought

[Question]. Think step by step.

04 Iterative

Tone -> format -> length -> constraints.

05 Delimiters

Wrap pasted content in "" or <tags>.

06 Instructional

Start with a verb. Summarize. Write. List.

07 Negative

Do not [X]. Do not [Y]. Do not [Z].

08 Chaining

Step 1 -> Step 2 -> Step 3.

09 RAG

Answer only from text below.
""[source]""

10 Emotion

This is important for my [career].

// WANT MORE

New Techniques, Every Week.

This cheatsheet covers the foundations. New techniques, AI tool breakdowns, and weekly creator retrospectives go out in the AIToolsDaily24 newsletter. Free. No spam. Built for people who actually use AI for work.

NEWSLETTER

Weekly AI techniques + tool breakdowns

longvy24.github.io/AI_Blog/index.html**YOUTUBE**

Full video versions of every technique

youtube.com/@AIToolsDaily24**TIKTOK**

60-second versions, posted daily

[@dj.bomasterflex](https://tiktok.com/@dj.bomasterflex)**INSTAGRAM**

Reels + carousels you can save

[@aitoolsdaily245](https://instagram.com/aitoolsdaily245)**BLOG**

Long-form breakdowns of every video

longvy24.github.io/AI_Blog/index.html**BUILT WITH PURPOSE**

This cheatsheet was assembled by Ryu Vy from real day-to-day AI use, with content drafted in collaboration with Claude and edited for accuracy. Share it freely with anyone learning to prompt.